




Bruno Stefoni

New York, NY | brunostefoni12@gmail.com | +1 646 203 7124 |  LinkedIn |  GitHub |  Scholar

PROFESSIONAL SUMMARY

Six years of experience with AI, Machine Learning and Data: deploying AI and ML models to the cloud; developing custom solutions with LLMs; building data pipelines for KPI tracking; designing A/B tests.

EXPERIENCE

- CloudWerx** Nov 2023 – Present
AI / ML Engineer Manhattan, New York
 - Project examples: RAG Chatbot; Augmented Search with LLM; fine-tuning pipeline for diffusion model.
 - Developing ML and AI solutions from client requirements using OSS, GCP, Vertex AI, and LLMs.
 - Consulting clients from: cyber security, healthcare, finance, ecommerce, advertising, robotics, architecture.
- Aetna - CVS Health** Sep 2023 – Oct 2023
Data Scientist Manhattan, New York
 - Designed A/B test experiment with 9 MM MAU App to measure campaign impact on customer adoption KPIs.
 - Analyzed data in GCP Big Query to perform demographic analysis for a live A/A experiment.
- Sortile** Jun 2023 – Aug 2023
Data Scientist Manhattan, New York
 - Built a 92% accurate XGBoost classifier to identify black/white fibers using light wavelength sensor data.
 - Developed and benchmarked XGBoost, LGBM and Deep Learning models. Deployed XGB model on AWS Lambda.
- Google** Jun 2022 – Aug 2022
Data Scientist Intern Mountain View, California
 - Conducted detailed market research and graph analysis by building a network dataset from diverse data sources using PLX, BigQuery and Python.
 - Detected up to 15% of user churn by continuous survival and time series analysis of Software adoption rate data.
 - Presented insights to internal stakeholders and management. Proposed a new methodology for analysts.
- STP** May 2018 – May 2022
Data Scientist Santiago, Chile
 - Built and implemented a company-wide chat bot on prem cloud server to support management and on-site workers.
 - Drove STP's revenue by identifying routes with up to 25% difference between planned and real bus speeds.
 - Prevented \$500K investment on less accurate tech by launching a statistical study.
 - Detected data outliers due to faulty batteries in 20% of STP's electric fleet via statistical analysis.
 - Increased main contract KPI by 6% by scraping Twitter data to alert service disruptions.

PROJECTS AND PUBLICATIONS

- Anky** Developed RAG + SFT/DPO fine tuned LLM with llama.cpp using social network data
- S&P Global Capstone Project** News classification by fine-tuning BERT from unlabeled data (97% accuracy)
- Master's Thesis Publications** Two papers published on computer microscopic simulation
- Deep Learning Competition** Next frame video prediction using Transformer and CNN (ResNet) in JEPA network
- Text Topic Analysis** LDA and GloVe word embeddings on large academic thesis corpus
- Twitter Sentiment Analysis** Using fastText and fine-tuning BERT for content strategy assessment
- Hackathon Winner** Led team of 6, designed and developed an app for transit users

EDUCATION

- NEW YORK UNIVERSITY** 2021 – 2023
Master of Science in Data Science GPA: 3.8/4.0
 - Coursework: Deep Learning, Machine Learning, Big Data, Optimization & Linear Algebra, Probability & Statistics
 - TA of Data Science Class at NYU Business School for two semesters

•UNIVERSIDAD DE CHILE

Bachelor and Master's degree in Engineering

2018

Final grade 6.6/7.0

- Graduated with highest distinction in top 10% GPA of cohort
- Granted Outstanding Student Award
- TA of five different courses, e.g., Stochastic Processes, Network Optimization
- Key Modules: Data Mining, Computer Science I and II, Probability, Statistics, Optimization, Operations Research

SKILLS

Certifications: Google Professional Machine Learning Engineer

Programming Languages: Python, R, MATLAB, Java

AI Libraries: Transformers, Langchain, OpenAI API, Anthropic API, Pydantic, HuggingFace

ML Libraries: PyTorch, PyTorch-Lightning, XGBoost, LGBM, scikit-learn, statsmodels

Cloud Platforms: Google Cloud Platform (GCP), AWS

Databases: SQL (SQLite), PostgreSQL (AlloyDB), BigQuery, Vector DBs (Milvus, Pinecone)

DevOps: Git, Docker, Terraform, Bash, Linux

Data Analysis: Pandas, NumPy, Plotly, Seaborn, Matplotlib

NLP: NLTK, SpaCy, fastText

Big Data: pySpark, Dask

Languages: Native English and Spanish speaker, Intermediate Italian